

Lecture 4 — Learning to Price: Optimism and UCB

Lecturer: Sid Banerjee

Date: September 3, 2026

4.1 Overview

Lecture 3 introduced the stochastic multi-armed bandit problem and asked:

How close can we get to a decision maker who knows the reward means in advance?

The key difficulty was that the data we observe depend on the actions we choose. Purely greedy decision making can lock onto the wrong arm, while uniform exploration wastes samples on arms that are already clearly inferior.

To design an adaptive policy, we first need to quantify what an empirical mean does—and does not—tell us after finitely many observations.

We now develop a more adaptive principle:

When uncertain, act as if the world might be favorable.

This is *optimism under uncertainty*. It leads to the Upper Confidence Bound (UCB) algorithm. The mathematical goal is to understand why a suboptimal arm with gap Δ_i is selected only about

$$O\left(\frac{\log T}{\Delta_i^2}\right)$$

times, leading to regret

$$O\left(\sum_{i:\Delta_i>0} \frac{\log T}{\Delta_i}\right).$$

We will then return to pricing and discuss what is special about price as a bandit action. Here is the proof structure we will use:

empirical mean $\longrightarrow$ confidence interval $\longrightarrow$ UCB $\longrightarrow \Delta_i \leq 2r_i \longrightarrow$ regret.

Sections marked * are optional reading.

4.2 Empirical means and finite-sample confidence

Recall that there are K arms. Arm i has iid rewards in $[0, 1]$ with unknown mean μ_i . Imagine that arm i has an infinite reward tape

$$X_{i,1}, X_{i,2}, \dots, \quad \mathbb{E}[X_{i,s}] = \mu_i.$$

After observing the first n rewards from this arm, its empirical mean is

$$\widehat{M}_{i,n} = \frac{1}{n} \sum_{s=1}^n X_{i,s}. \quad (4.1)$$

The capital M emphasizes that $\widehat{M}_{i,n}$ is random before the data are observed; after observing a particular history, its realized value may be written $\widehat{m}_{i,n}$. The population mean μ_i is fixed but unknown.

Let

$$\mu^* = \max_i \mu_i, \quad \Delta_i = \mu^* - \mu_i.$$

Let

$$N_i(t) = \sum_{s=1}^t \mathbf{1}\{I_s = i\}$$

be the number of times arm i is selected in the first t rounds, and let $\widehat{M}_i(t) = \widehat{M}_{i,N_i(t)}$ be the empirical mean based on the rewards revealed from that arm by time t .

As in Lecture 3,

$$\text{Reg}(T) = \sum_i \Delta_i \mathbb{E}[N_i(T)].$$

Thus the central issue is to control how often suboptimal arms are played.

For a fixed arm i and a fixed sample size n , Hoeffding's inequality gives

$$\mathbb{P}\left(|\widehat{M}_{i,n} - \mu_i| \geq r\right) \leq 2e^{-2nr^2}.$$

Choose

$$r(n) = \sqrt{\frac{2 \log T}{n}}. \quad (4.2)$$

Then

$$\mathbb{P}\left(|\widehat{M}_{i,n} - \mu_i| > r(n)\right) \leq 2T^{-4}.$$

The important scale is

$$\boxed{r(n) \asymp \frac{1}{\sqrt{n}}}.$$

Four times as many observations are needed to cut statistical uncertainty in half. The logarithmic factor makes the intervals reliable simultaneously over many arms and sample sizes.

4.2.1 A clean event

We would like all confidence intervals to be valid simultaneously over all arms and all sample sizes. Define

$$\mathcal{E} = \left\{ |\widehat{M}_{i,n} - \mu_i| \leq r(n) \text{ for every } i \in [K], n \leq T \right\}. \quad (4.3)$$

A union bound gives

$$\mathbb{P}(\mathcal{E}^c) \leq 2KT T^{-4}.$$

If $K \leq T$, this is at most $2/T^2$.

On the clean event, we can stop doing probability and reason deterministically using the confidence intervals.

Remark 4.1 (Why sample size is subtle). *In an adaptive algorithm, $N_i(t)$ is random and depends on previous rewards. The reward-tape construction above separates this adaptivity from concentration: the n th time arm i is played, the algorithm simply reveals $X_{i,n}$. Hoeffding is applied to each fixed prefix length n , and then union-bounded over n and i . This is the elementary device used in Slivkins' treatment of stochastic bandits.*

4.3 Upper confidence bounds

Define

$$\text{UCB}_i(t) = \widehat{M}_i(t-1) + r(N_i(t-1)), \quad (4.4)$$

$$\text{LCB}_i(t) = \widehat{M}_i(t-1) - r(N_i(t-1)). \quad (4.5)$$

The UCB algorithm is:

1. Play each arm once.
2. For $t > K$, choose

$$I_t \in \arg \max_i \text{UCB}_i(t). \quad (4.6)$$

Why is this sensible? An arm has a large UCB for one of two reasons:

- its empirical mean is large — a reason to *exploit*;
- its confidence radius is large — a reason to *explore*.

UCB combines these two motives in a single index.

Example 4.2 (A two-interval snapshot). Suppose the current confidence intervals are

$$[0.48, 0.62] \quad \text{and} \quad [0.35, 0.75].$$

The second arm has the lower empirical center, but the larger upper endpoint. UCB selects it because the available data have not yet ruled out that it is best. Another sample is valuable precisely because this uncertainty could change the decision.

4.4 The key UCB inequality

Let i^* be an optimal arm. Suppose UCB chooses arm i at round t . By definition,

$$\text{UCB}_i(t) \geq \text{UCB}_{i^*}(t).$$

On the clean event,

$$\text{UCB}_{i^*}(t) \geq \mu^*.$$

Also,

$$\widehat{M}_i(t-1) \leq \mu_i + r(N_i(t-1)),$$

so

$$\text{UCB}_i(t) = \widehat{M}_i(t-1) + r(N_i(t-1)) \leq \mu_i + 2r(N_i(t-1)).$$

Therefore

$$\mu^* \leq \mu_i + 2r(N_i(t-1)), \tag{4.7}$$

or

$$\boxed{\Delta_i \leq 2r(N_i(t-1))}. \tag{4.8}$$

This is the whole proof in one line:

If UCB still chooses a bad arm, then that arm must still be sufficiently uncertain.

Substituting (4.2) into (4.8),

$$\Delta_i \leq 2\sqrt{\frac{2 \log T}{N_i(t-1)}}.$$

Rearranging,

$$N_i(t-1) \leq \frac{8 \log T}{\Delta_i^2}. \tag{4.9}$$

Thus, apart from initialization and the rare bad event, a suboptimal arm cannot be selected after it has accumulated more than order $\log T / \Delta_i^2$ samples. Its final count can exceed the displayed threshold by at most the last selected round.

4.5 Gap-dependent regret

Using the regret decomposition,

$$\text{Reg}(T) = \sum_{i:\Delta_i>0} \Delta_i \mathbb{E}[N_i(T)].$$

Equation (4.9) suggests

$$\mathbb{E}[N_i(T)] = O\left(1 + \frac{\log T}{\Delta_i^2}\right).$$

Hence

$$\boxed{\text{Reg}(T) = O\left(\sum_{i:\Delta_i>0} \frac{\log T}{\Delta_i}\right)}. \quad (4.10)$$

The contribution of the bad event is negligible because regret is at most T while $\mathbb{P}(\mathcal{E}^c) = O(T^{-2})$.

Remark 4.3 (Instance-dependent versus worst-case guarantees). The bound (4.10) is excellent when every suboptimal arm is separated from the optimum by a nontrivial gap. But it becomes large when some Δ_i is very small. A complementary argument splits arms into “small-gap” and “large-gap” groups and yields an instance-independent bound of order

$$O\left(\sqrt{KT \log T}\right).$$

Sharp constants are secondary here; the important point is the qualitative difference between logarithmic instance-dependent regret and square-root worst-case regret.

4.6 Optional: successive elimination *

UCB is not the only way to use confidence intervals. Another natural algorithm repeatedly samples all active arms and removes arm i once

$$\text{UCB}_i(t) < \text{LCB}_j(t)$$

for some other active arm j .

The same clean-event proof implies that the best arm is never eliminated, and an arm with gap Δ_i is discarded after roughly

$$O\left(\frac{\log T}{\Delta_i^2}\right)$$

samples.

This alternate view is useful because it isolates the deeper principle:

adaptive exploration means spending samples only until an action is statistically resolved.

4.7 Optional: why exploration is unavoidable ^{*}

Could a clever algorithm avoid trying uncertain arms altogether? No. To know whether an arm with mean $1/2$ is actually slightly better or slightly worse than another arm, one needs enough samples to distinguish nearby probability distributions.

For a Bernoulli coin with mean

$$\frac{1}{2} + \delta$$

versus one with mean

$$\frac{1}{2} - \delta,$$

the standard error after n samples is of order $1/\sqrt{n}$. Distinguishing the two therefore requires roughly

$$n = \Omega\left(\frac{1}{\delta^2}\right)$$

samples. More refined lower bounds use KL divergence, but the core message is already visible here:

learning a gap of size δ costs on the order of $1/\delta^2$ observations.

This explains why the UCB sample complexity in (4.9) has exactly that dependence, up to the logarithmic confidence factor.

4.8 Back to pricing: prices as arms

Suppose a seller chooses from a finite price set

$$\mathcal{P} = \{p_1, \dots, p_K\}.$$

At round t , price p_i is offered. Let

$$Y_t = \mathbf{1}\{\text{customer purchases at round } t\}.$$

Conditional on choosing price p_i ,

$$Y_t \sim \text{Bernoulli}(q_i), \quad q_i = 1 - F(p_i),$$

where q_i is unknown to the seller. The realized revenue and its conditional mean are

$$X_t = p_i Y_t, \quad \mu_i = p_i q_i.$$

After normalizing by the largest possible price, this is exactly a stochastic bandit.

UCB is a canonical model for the exploration–exploitation phenomenon. It should not be read as a claim that a production pricing system literally treats every price as an unrelated arm; the structured feedback discussed below often supports more specialized algorithms.

So a direct pricing implementation of UCB is simply:

$$\boxed{\text{offer the price with the largest optimistic estimate of expected revenue.}} \quad (4.11)$$

This captures the idea that

$$\boxed{\text{a price is both a revenue decision and an experiment.}}$$

4.9 But pricing has extra structure

Treating each candidate price as an unrelated arm throws away useful information.

Suppose a customer has a fixed valuation V . If we offer price p and observe a sale, then we learn

$$V \geq p.$$

This implies that the customer would also have purchased at every lower price. Conversely, if there is no sale at p , then

$$V < p,$$

so the customer would not have purchased at any higher price.

These are counterfactual statements about the same customer. Under an iid valuation model, they are also censored information about F , and hence about purchase probabilities at other prices. Thus pricing has *partial feedback* across prices:

- sale at p gives information about all lower prices;
- no sale at p gives information about all higher prices.

There is also functional structure. The mean revenue is not an arbitrary vector:

$$\mu(p) = p(1 - F(p)).$$

Nearby prices typically have related purchase probabilities and revenues.

The optional binary-logit model in Lecture 3 imposes still more structure:

$$q(p, x) = \sigma(\theta^\top x - \beta p).$$

One observation can then update common parameters and inform many prices at once. This suggests a useful hierarchy:

$$\boxed{\text{unstructured finite arms} \longrightarrow \text{parametric demand} \longrightarrow \text{ordered or continuum prices.}}$$

Moving right can improve learning by sharing information, but only when the additional modeling assumptions are credible.

The basic UCB reduction ignores both forms of structure. More specialized dynamic-pricing algorithms exploit them.

4.10 Optional: a glimpse of continuum pricing ^{*}

Suppose prices can be any value in an interval, say $p \in [0, 1]$. One simple idea is to discretize into K grid points and run a finite-armed bandit algorithm.

This creates two sources of regret:

1. **learning error:** regret from not knowing which grid price is best;
2. **discretization error:** regret because the true optimal price may lie between grid points.

If the revenue curve is sufficiently smooth, a finer grid reduces discretization error but creates more arms to learn. Choosing the grid therefore involves a bias–variance-like tradeoff.

This is the starting point of the dynamic-pricing literature exemplified by Kleinberg–Leighton: exploit the order and smoothness of prices rather than treating prices as unrelated actions. We will not develop the full theory here; the finite-armed model already gives us the core exploration–exploitation ideas that recur later in the course.

4.11 Proof pattern to remember

The UCB proof illustrates a template we will use again:

1. define an ideal benchmark;
2. construct an empirical estimate and a high-probability “clean event”;
3. show that a wrong decision on the clean event can happen only while some uncertainty remains large;
4. bound the number of such decisions;
5. charge the rare complement of the clean event separately.

More rounds shrink uncertainty only when the decision rule continues to generate informative observations. UCB is one way to convert scale into learning: uncertainty itself forces the samples needed to resolve it.

Later, the Bayes Selector analysis will use a different algorithm and benchmark, but a related high-level idea: control regret by controlling the probability or frequency with which the online action disagrees with an ideal offline decision.

4.12 Takeaways

1. Greedy fails because uncertainty and low quality are confounded.
2. Confidence intervals quantify uncertainty.
3. UCB uses optimism to turn uncertainty into exploration.
4. A suboptimal arm with gap Δ needs only about $O(\log T/\Delta^2)$ samples before it is resolved.
5. This yields logarithmic instance-dependent regret.
6. Dynamic pricing is a bandit problem, but with useful order, smoothness, and partial-feedback structure across prices.

Recommended reading

Core / closest to this lecture.

Slivkins, Chapter 1, Sections 1.3.1–1.3.3. These sections develop the reward-tape clean event, confidence bounds, successive elimination, and optimism/UCB in essentially the proof style used here.

Alternative perspective.

Re-read Lecture 3's pricing formulation while translating every arm i into the action "charge price p_i ." This makes clear why the generic UCB theorem applies after revenue normalization, and what economic information the abstract arm model suppresses.

Optional / deeper.

Slivkins, Chapter 2, gives lower bounds showing that exploration is unavoidable. The earlier ORIE 6154 notes on bandits in revenue management and the Kleinberg–Leighton dynamic-pricing framework study pricing-specific extensions that exploit order, smoothness, and feedback shared across prices; their technical proofs are not required here.